UNIVERSITY OF COPENHAGEN
FACULTY OF SCIENCE

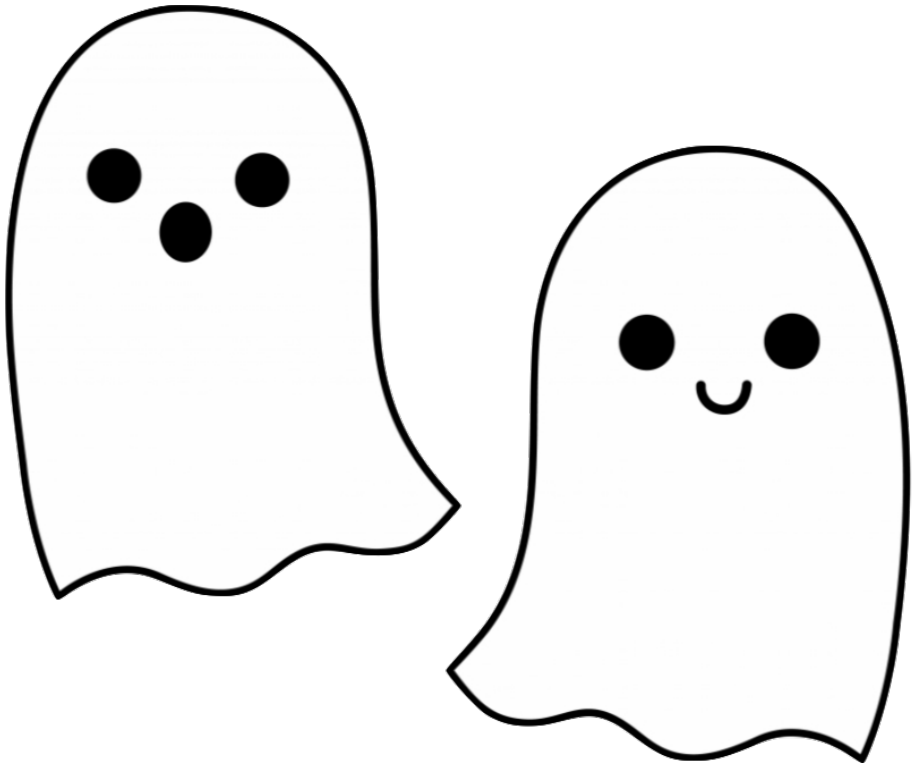

Project outside course scope

Jacob Herbst (mwr148), Matilde Broløs (jtw868)

IFC

An application for dynamic evaluation and static verification of programs

Advisor: Ken Friis Larsen

1st of April 2022

Abstract

Contents

1 Introduction **5**

2 Background **6**
2.1 Language . . . 6
2.2 Hoare Logic . . . 9
2.2.1 Assertions . . . 10
2.2.2 Axiomatic system for partial correctness . . . 10
2.2.3 Total correctness . . . 11
2.2.4 Soundness and Completeness . . . 12
2.3 Verification Condition Generation . . . 12
2.3.1 Weakest Liberal Precondition . . . 12
2.3.2 Weakest Precondition . . . 14
2.4 Verifying conditions with automated solvers . . . 15

3 Implementation **15**
3.1 Parser . . . 16
3.1.1 Grammar ambiguity . . . 16
3.1.2 Ghost variables . . . 17
3.2 Interpreter . . . 18
3.3 Verification Condition Generator . . . 19
3.3.1 First attempt . . . 20
3.3.2 Second attempt . . . 20
3.3.3 While loops with invariants and variants . . . 21
3.3.4 Handling undefined behaviour in Verification Condition Generator . . . 22
3.4 Proof-assistant API . . . 22
3.5 Interface for proofs . . . 23

4 Assessment **23**
4.1 Experiments . . . 24
4.1.1 Blackbox testing . . . 25
4.1.2 Quickchecking instances . . . 25
4.1.3 Dynamic execution compared to static proofs . . . 26
4.2 Evaluation . . . 29
4.2.1 Correctness . . . 30
4.2.2 Completeness . . . 30
4.2.3 Robustness . . . 30
4.2.4 Maintainability . . . 31

5 Discussion and conclusion **31**
5.1 Future work . . . 32
5.1.1 Procedures and arrays . . . 32
5.1.2 Type system . . . 32
5.1.3 PER-logic for Information Flow Control . . . 32
5.2 Conclusion . . . 33

A How To Use **35**

B AST **36**

C Table of test programs **37**

D Example programs **40**
D.1 `always_wrong.ifc` 40
D.2 `assign.ifc` 40
D.3 `collatz.ifc` 40
D.4 `div.ifc` 40
D.5 `div_in_cond.ifc` 40
D.6 `fac.ifc` 41
D.7 `fakesum.ifc` 41
D.8 `false_mult.ifc` 41
D.9 `fib.ifc` 41
D.10 `infinity.ifc` 42
D.11 `isqrt.ifc` 42
D.12 `isqrt_fast.ifc` 42
D.13 `isqrt_sub.ifc` 43
D.14 `max.ifc` 43
D.15 `mccarthy.ifc` 43
D.16 `mod0.ifc` 44
D.17 `mult.ifc` 44
D.18 `mult3.ifc` 44
D.19 `skip.ifc` 44
D.20 `sum.ifc` 45
D.21 `undef.ifc` 45
D.22 `var.ifc` 45

1 Introduction

In this project we present a small imperative *While* language in the C-style family with a built-in assertion language. This language uses Hoare logic and predicate tranformer semantics to allow for verification condition generation similar to Why3. Our VC generator discharges verification conditions fit for the SMT solver Z3, allowing us to prove the partial or total correctness of a program.
The main motivation for implementing this language is to make a base which we can easily extend to include *Partial Equivalence Relation* (PER) logic, which is a logic proposed to reason about *Information Flow Control*[4]. This can be used for making assertions about the security of code.
Predicate Transformer Semantics, which is used for verification condition generation, was introduced by Dijkstra in 1975[3]. These conditions are based on the work of Robert W. Floyd, who presented a form of program verification using logical assertions in his paper from 1967[5]. Floyds work led to the axiomatic system presented by Sir C.A.R Hoare in 1969[6], also known as Floyd-Hoare logic. Using this system, and the rules of inference that it introduces, one can formally prove the correctness of a program. Combined, this provides an approach for automatically generating conditions which can verify the correctness of certain programs.
Why3 is an example of a platform which provides automatic program verification. Why3 requires the programmer to write assertions inside the program code, and then utilises these assertions to compute verification conditions for the program. Once the verification condition is generated, Why3 can discharge this condition to a variety of SMT solvers or proof assistants, which can then determine whether this condition can be proved.
This report presents the work of this project as follows:

- In Section 2, we present the syntax and semantics of our language. Next we explain how VC generation works, how it is related to Hoare logic and how it can be used for automatically proving certains properties about programs. Finally we describe the coupling between VC-generators and SMT-solvers.
- In Section 3, we present our implementation, which can both dynamically evaluate a program and statically prove the correctness. In the implementation we transform the formal semantics of the language into equivalent Haskell code and utilize predicate transformer semantics to compute verification conditions for a program.
- In Section 4, we present our strategy for evaluting the code quality and give an assessment of the implementation. Here, a combination of blackbox tests, automatic tests, and test programs, is used to assess the quality.
- Lastly, Section 5 presents some ideas for future work, gives a discussion of the work, and concludes the project.

Throughout the project our main focus has been on getting a good understanding of how verification conditions work, and how to convert this into Haskell code. It has also been on building an implementation that is robust and easily extended and maintained. Furthermore it has been interesting for us to investigate what is necessary to prove correctness of a program, using our implementation, and how to thouroughly test the correctness of our work.
Enjoy!

2 Background

In this section, we present the grammar and semantics of the *While* language used throughout the project. Next, we explain Floyd-Hoare logic and verification condition generation through semantic predicate transformers and how this provides a basis for formal verification. Finally, we describe how this can be coupled with SMT solvers to allow automatic proving of programs.

2.1 Language

While is a small imperative programming language with a built-in assertion language, enabling the possibility of statically verifying programs by the use of external provers. The assertions also enable us to dynamically check the program during evaluation. This section presents the syntax and semantics of both the imperative language and the assertion language.
The language is simple, as the focus in this project has been to correctly generate verification conditions for said programs. Table 1 shows the grammar of *While*. In essence, a *While* program is a statement with syntax in the C-family. Although not strictly part of the language itself, we require a program to have a *header*, succeeded by one or more statements. We leave out the syntax of the header in the grammar but will describe it in detail in Section 3.5, where we explain why it is necessary as well.
The program syntax is straightforward and does, in general, not provide anything new. One of the interesting program constructs are while loops. A while loop consists of a loop condition, one or more invariants, and zero or one variant, followed by a loop body. The reason for the inclusion of loop invariants and variants will become apparent in Section 2.2. Another interesting construct is if statements, which can be constructed both with and without an else branch. Finally, assertion statements are of great importance for this project. An assertion is denoted by a hashtag and curly braces. Assertions are explained in detail in Section 2.2 and Section 2.3.

The general syntax of the language is as follows. Identifiers must start either with a lowercase symbol or an underscore, and may then be proceeded by 0 or more characters, digits or underscores. Ghost variables must start with either a *ghost*-emoji or, for users more comfortable with ascii, a $, directly succeeded by an identifier.

Integers can be specified as decimal, hexadecimal, binary and octal.

While includes many common operators, for which the syntax follows the usual conventions. For arithmetic expressions (`aexpr`) we have the following precedence:

- Parentheses
- Negation
- Multiplication, division and modulo (left associative)
- Addition and subtraction (left associative)

where parentheses binds tightest. Operators on boolean expressions (`bexpr`) follow the standard precedence rules:

- !

⟨*statements*⟩ ::= ⟨*statement*⟩ ';' ⟨*statements*⟩ | ϵ

⟨*statement*⟩ ::= ⟨*ghostid*⟩ ':=' ⟨*aexpr*⟩
| ⟨*identifier*⟩ ':=' ⟨*aexpr*⟩
| 'if' ⟨*bexpr*⟩ '{' ⟨*statements*⟩ '}'
| 'if' ⟨*bexpr*⟩ '{' ⟨*statements*⟩ '}' 'else' '{' ⟨*statements*⟩ '}'
| 'while' ⟨*bexpr*⟩ ⟨*invariant*⟩ ⟨*variant*⟩ '{' ⟨*statements*⟩ '}'
| '#{' ⟨*assertion*⟩ '}'
| 'skip' | 'violate'

⟨*invariant*⟩ ::= '?{' ⟨*assertion*⟩ '}'
| ⟨*invariant*⟩ ';' ⟨*invariant*⟩

⟨*variant*⟩ ::= '!{' ⟨*aexpr*⟩ '}' | ϵ

⟨*assertion*⟩ ::= '$\forall$' ⟨*identifier*⟩ '.' ⟨*assertion*⟩
| '$\exists$' ⟨*identifier*⟩ '.' ⟨*assertion*⟩
| '$\sim$' ⟨*assertion*⟩
| ⟨*assertion*⟩ ⟨*assertionop*⟩ ⟨*assertion*⟩
| ⟨*bexpr*⟩

⟨*assertionop*⟩ ::= '$\wedge$' | '$\vee$' | '$\Rightarrow$'

⟨*bexpr*⟩ ::= 'true' | 'false' |
| '!'⟨*bexpr*⟩
| ⟨*bexpr*⟩ ⟨*bop*⟩ ⟨*bexpr*⟩
| ⟨*bexpr*⟩ ⟨*rop*⟩ ⟨*bexpr*⟩
| '(' ⟨*bexpr*⟩ ')'

⟨*bop*⟩ ::= '&&' | '||'

⟨*rop*⟩ ::= '<' | '$\leq$' | '=' | '$\neq$' | '>' | '$\geq$'

⟨*aexpr*⟩ ::= ⟨*identifier*⟩
| ⟨*ghostid*⟩
| ⟨*integer*⟩
| '-'⟨*axpr*⟩
| ⟨*aexpr*⟩ ⟨*aop*⟩ ⟨*aexpr*⟩
| '(' ⟨*aexpr*⟩ ')'

⟨*aop*⟩ ::= '$+$' | '$-$' | '$*$' | '/' | '%'

⟨*ghostid*⟩ ::= ⟨*ghost emoji*⟩ ⟨*identifier*⟩ | \$⟨*identifier*⟩

Table 1: Grammar of IFC

- && (right associative)
- || (right associative)

and the *First Order Logic* operators used in assertions (`assertion`) follow the precedence rules:

- $\sim$
- $\wedge$ (right associative)
- $\vee$ (right associative)
- $\forall$ and $\exists$
- $\Rightarrow$ (right associative)

where $\sim$ binds tightest. Relational operators works on *aexpr* and are non-associative.

An important point about the syntax is that we want to make a clear distinction between boolean expressions in the *While* language and the *First Order Logic* used in assertions. Thus, only one of the syntaxes for conjunction and disjunction is allowed, based on the context. We find that this makes the difference between assertions and program logic more evident. Furthermore, it eliminates any complications with using the same operator symbol for different program constructs, which could cause ambiguity.
In Table 2 we show the semantics of the different statements in *While*. The program variables are kept in a store, mapping variables to integers. If undefined behavior occurs, we have an abnormal termination. Thus the evaluation of commands goes from a store to a (potentially new) store or an abnormal termination, denoted by $\lightning$.
We have left out the semantics for arithmetic and boolean expressions, as these follow the standard semantics. However, it should be noted that we do not allow *division* and *modulo* by zero. Thus, if one of these occurs, the program will terminate with an abnormal store. Furthermore, ghost variables can only occur in assertions and not as part of the program logic. We will come back to this in Section 2.2.
The semantics describe how to evaluate statements in *While*. The most interesting semantic rules in our language are the ones for while loops. Here we have to assert that the loop invariant holds both at the beginning of each loop iteration and after loop termination. If the invariant ever evaluates to *false*, the program terminates abnormally, as it should for all assertions.

```
1 res := 0;
2 while (q > 0) {
3       res := res + r;
4       q := q - 1;
5 };
```

Figure 1: Example program `mult.ifc`

To show what the language is capable of, a small example program computing the multiplication of two integers q and r is presented in Figure 1. The program takes as input two integers q and r, and multiplies them using a naive approach utilising a while loop. We will use this program as an example throughout the report. Line 1 shows the syntax

$violate$	$\frac{}{\langle \texttt{violate}, \sigma \rangle \downarrow \lightning}$
$skip$	$\frac{}{\langle \texttt{skip}, \sigma \rangle \downarrow \sigma}$
$assign$	$\frac{\langle a, \sigma \rangle \downarrow n}{\langle X := a, \sigma \rangle \downarrow \sigma[X \mapsto n]}$
$sequence$	$\frac{\langle s_0, \sigma \rangle \downarrow \sigma'' \quad \langle s_1, \sigma'' \rangle \downarrow \sigma'}{\langle s_0 ; s_1, \sigma \rangle \downarrow \sigma'}$
$if\text{-}true$	$\frac{\langle b, \sigma \rangle \downarrow \texttt{true} \quad \langle s_0, \sigma \rangle \downarrow \sigma'}{\langle \texttt{if } b \texttt{ then } s_0 \texttt{ else } s_1 \rangle \downarrow \sigma'}$
$if\text{-}false$	$\frac{\langle b, \sigma \rangle \downarrow \texttt{false} \quad \langle s_1, \sigma \rangle \downarrow \sigma'}{\langle \texttt{if } b \texttt{ then } s_0 \texttt{ else } s_1 \rangle \downarrow \sigma'}$
$while\text{-}false$	$\frac{\langle b, \sigma \rangle \downarrow \texttt{false} \quad \langle i, \sigma \rangle \downarrow \texttt{true}}{\langle \texttt{while } b \texttt{ invariant } i \texttt{ do } s_0 \rangle \downarrow \sigma}$
$while\text{-}true$	$\frac{\langle b, \sigma \rangle \downarrow \texttt{true} \quad \langle i, \sigma \rangle \downarrow \texttt{true} \quad \langle s_0, \sigma \rangle \downarrow \sigma'' \quad \langle \texttt{while } b \texttt{ invariant } i \texttt{ do } s_0, \sigma'' \rangle \downarrow \sigma'}{\langle \texttt{while } b \texttt{ invariant } i \texttt{ do } s_0 \rangle \downarrow \sigma'}$
$while\text{-}i\text{-}false$	$\frac{\langle i, \sigma \rangle \downarrow \texttt{false}}{\langle \texttt{while } b \texttt{ invariant } i \texttt{ do } s_0 \rangle \downarrow \lightning}$

Table 2: Semantics for the *While* language.

of assignments, lines 2-5 shows the syntax of while loops, and the entire program is also a demonstration of sequencing statements. For now, we leave out the assertions, but we will come back to this in Section 2.2.

Even though this small *While* language is very simple, it is still very interesting as the base language for our project. By creating an assertion language that can prove the correctness of programs written in this simple language, we can investigate the possibility of proving termination and correctness of while loops, which are definitely the tricky construct of this language. Loops are specifically challenging because we cannot know whether they ever terminate. Furthermore, the way a loop affects program variables is less see-through than, for example, the effects of an assignment. Therefore, it is desirable to prove the correctness, and also termination if possible, of such program constructs. That is why this small language is well suited for this project.

2.2 Hoare Logic

Now we look at Hoare logic and how we can use it to verify our program. To assert that a program works as intended, we want to be able to prove it formally. To avoid the proofs being too detailed and comprehensive, one wants to look at the essential properties of the constructs of the program. We look at partial correctness of a program, meaning that we do not ensure that a program terminates, only that *if* it terminates, certain properties will hold.

```
1  vars: [q,r]
2  requirements: {q >= 0}
3  <!=_=!>
4  res := 0;
5  $a := q;
6  while (q > 0) ?{res = ($a - q) * r /\ q >= 0} !{q}{
7    res := res + r;
8    q := q - 1;
9  };
10 #{res = $a * r };
```

Figure 2: Example program `mult.ifc`

2.2.1 Assertions

For expressing properties of a program, we use *assertions*, which are a first order logic formula, which should hold at a specific time in a program. By combining these into a triple with a *precondition* P and a *postcondition* Q and a statement S we get:

$$\{P\}S\{Q\}$$

This triple states that *if* the precondition P holds in the initial state, and *if* S terminates when executed from that state, *then* Q will hold in the state in which S halts. Thus the triple does not say anything about whether S terminates, only that if it terminates, we know Q to hold afterward. Moreover, if S does not terminate, then any Q will hold. These triples are called Hoare triples.

Logical variables vs program variables. When using assertions we differ between *program variables* and *logical variables*. Program variables are usable in the program logic and are mutable throughout the program. Sometimes we might need to keep the original value of a variable that has changed during the program. Here we can use a *logical variable*, or *ghost variable*, to maintain the value. The ghost variables can only be used in assertions, not in the actual program, and thus cannot be changed.

Multiplication example. To show how we can use assertions, we look at the multiplication program presented before. The version with assertions is shown in Figure 2.

In line 1-2, we have the input variables listed, and the requirements for the input given. This is part of the program header, which we will describe in more detail later. The requirements state that q must be greater than or equal to zero, since we would otherwise never get a correct result, as a negative value for q would prevent us from going into the while loop. Clearly the result is correct if q is zero.
Additional logic using if-else statements would alleviate this requirement for a complete function that would also work with negative values for q. In line 10, we assert that the final result will be the original value of q multiplied by r. Here we use a ghost variable for keeping the original value of q. We also have an invariant and a variant in the while loop in line 6, but we will come back to the meaning of that in Section 2.2.3.

2.2.2 Axiomatic system for partial correctness

Hoare logic specifies an inference system for partial correctness of a program based on the semantics of the different constructs of the language. The axiomatic system is

shown in Table 3.
This axiomatic system shows how assertions are evaluated in the Hoare logic. For example, for an assignment, we can say that if we bind x to the evaluated value of a in the initial state, and if P holds in this state, then after assigning x to a, P must still hold. For the *skip* command, we see that the assertion must hold both before and after, as *skip* does nothing. The axiomatic semantics for an assertion around a sequence of statements $\{P\}S_1; S_2\{R\}$ state that if P holds in the initial state, and executing S_1 in this state produces a new state in which Q holds, and if executing S_2 in this new state produces a state in which R holds, then $\{P\}S_1; S_2\{R\}$ holds.
Another interesting construct is the while loop. Here P denotes the loop invariant, and b is the loop condition. If b evaluates to $true$ and P holds in the initial environment, and executing S in this environment produces a new state in which P holds, then we know that after the while loop has terminated, we must have a state where P holds and where b evaluates to $false$.

$[violate_p]$	$\{false\}\texttt{violate}\{Q\}$
$[skip_p]$	$\{P\}\texttt{skip}\{P\}$
$[ass_p]$	$\{P[x \mapsto \mathcal{A}[\![a]\!]]\}\ x := a\ \{P\}$
$[seq_p]$	$\frac{\{P\}S_1\{Q\} \quad \{Q\}S_2\{R\}}{\{P\}S_1;S_2\{R\}}$
$[if_p]$	$\frac{\{\mathcal{B}[\![b]\!]\wedge P\}S_1\{Q\} \quad \{\neg\mathcal{B}[\![b]\!]\wedge P\}S_2\{Q\}}{\{P\}\ \texttt{if}\ b\ \texttt{then}\ S_1\ \texttt{else}\ S_2\{Q\}}$
$[while_p]$	$\frac{\{\mathcal{B}[\![b]\!]\wedge P\}S\{P\}}{\{P\}\ \texttt{while}\ b\ \texttt{do}\ S\{\neg\mathcal{B}[\![b]\!]\wedge P\}}$
$[cons_p]$	$\frac{\{P'\}S\{Q'\}}{\{P\}S\{Q\}}$ if $P \Rightarrow P'$ and $Q' \Rightarrow Q$

Table 3: Axiomatic system for partial correctness of *While*[8].

The syntax of our version of *While* deviates slightly from the syntax presented in Table 2 and Table 3. The reason why we have not depicted the exact syntax in the tables is that our use of curly brackets obscured the meaning of Hoare Triples, because of the conventional notation using curly brackets for indicating the pre- and postcondition. However, it is clear that the syntax presented in our grammar, and the syntax presented in the tables, are semantically equivalent.

2.2.3 Total correctness

The above states how the Hoare logic can prove partial correctness, but does not guarantee termination of loops. We want to be able to prove total correctness, meaning we want to prove termination as well. However, this cannot be done with the axiomatic system from Table 3, as the assertions are not sufficient to prove the termination of while loops. To verify that a program terminates, we need a more robust assertion in the form of a loop variant. The loop variant is an expression that needs to be subject to a well-founded relation. In our case, the only possible variant is an arithmetic expres-

sion that decreases with each loop iteration. For example, in the while loop from our example program q is the variant, as can be seen in line 8 of the code in Figure 2. We use the following modified syntax to express the logic of using a variant for a while loop.

$$\frac{\{I \wedge e \wedge v = \xi\} s \{I \wedge v \prec \xi\} \quad wf(\prec)}{\{I\} \texttt{ while } e \texttt{ invariant } I \texttt{ variant } v, \prec \texttt{ do } s \{I \wedge \neg e\}} \tag{1}$$

Where v is the variant of the loop, ξ is a fresh logical variable, and $\prec$ is a well-founded relation. Because the data type is unbounded integers, we use the following well-founded relation[7]:

$$x \prec y \quad = \quad x < y \wedge 0 \leq y$$

With both invariants and variants, it is possible to prove total correctness of programs containing while loops, as we can now reason about termination of loops.

2.2.4 Soundness and Completeness

For obvious reasons it is important that the axiomatic system is sound. This means that any property that can be proved by the inference system will be valid according to the semantics. We do not prove the soundness of the axiomatic system but simply refer to [8]. Hoare logic is however only shown to be relatively complete[2], which means that it is only as complete as the logical system of assertions. Hence we might have a program which is correct according to the semantics, but that we cannot prove to be correct.

2.3 Verification Condition Generation

In the previous section, we presented Hoare Triples and how they can be used to assign meaning to programs. As mentioned, for the postcondition to hold true, we want the program to be executed in a state satisfying the precondition. It is fair to assume that we know the postcondition of a program, but how can we know that our precondition is saying anything meaningful about our program? We can use *Weakest Precondition Calculus* (denoted WP) to generate a suitable precondition in such a case. It essentially states:

$$\{WP(S, Q)\} S \{Q\}$$

WP is a well-defined calculus to find the least restrictive precondition that will make Q hold after S. By this calculus, which we will present in this section, validation of Hoare Triples can be reduced to a single logical sentence. WP is functional compared to the relational nature of Hoare logic, which means that it is easier to reason about, as we dont have to consider any consequences. By computing the weakest precondition, we get a formula of first-order logic that can be used to verify the program. This process is called *verification condition generation*.

2.3.1 Weakest Liberal Precondition

The *Weakest Liberal Precondition* is the weakest precondition necessary for proving correctnes of a program, without proving termination. We denote it by WLP. In Figure 3, the structure for computing the WLP for the different constructs of *While* is shown.
Most of the rules are somewhat self-explanatory, but we would like to go through a selection of the rules that we find to be complex in our language setting.

$$
\begin{aligned}
WLP(\texttt{violate}, Q) &= false \\
WLP(\texttt{skip}, Q) &= Q \\
WLP(x := e, Q) &= \forall y, y = e \Rightarrow Q[x \leftarrow y] \\
WLP(s_1; s_2, Q) &= WLP(s_1, WLP(s_2, Q)) \\
WLP(\texttt{if } e \texttt{ then } s_1 \texttt{ else } s_2, Q) &= (e \Rightarrow WLP(s_1, Q)) \\
&\quad \wedge (\neg e \Rightarrow WLP(s_2, Q)) \\
WLP(\texttt{while } e \texttt{ invariant } I \texttt{ do } s, Q) &= I \wedge \\
&\quad \forall x_1, ..., x_k, \\
&\quad (((e \wedge I) \Rightarrow WLP(s, I)) \\
&\quad\quad \wedge ((\neg e \wedge I) \Rightarrow Q))[w_i \leftarrow x_i] \\
&\quad \text{where } w_1, ..., w_k \text{ is the set of} \\
&\quad \text{assigned variables in statement } s \text{ and} \\
&\quad x_1, ..., x_k \text{ are fresh logical variables.} \\
WLP(\{P\}, Q) &= P \wedge Q \quad \text{where P is an assertion}
\end{aligned}
$$

Figure 3: Rules for computing weakest liberal precondition of a statement in *While*[7].

Assignments. The rule for computing the WLP for assignments says that for all variables y where $y = e$, we should exchange x in Q with y. So we exchange all occurrences of x with the value assigned to x. We use a quantifier for y to avoid exponential growth by letting $y = e$ and substituting x with y in Q. For instance, if $e = 1 + 2 + ... + 100.000$, then by letting y hold the value of e, we avoid substituting in this very long expression, making the formula significantly more concise.

Sequence. For finding the *wlp* of a sequence of statements $s_1; s_2$, we need to first find the *wlp* Q' of s_2 with Q, and then compute the *wlp* of s_1 with Q'. This shows how we compute the weakest liberal precondition using a bottom-up approach.

While-loops. To compute the WLP of a while loop, we have to ensure that the loop invariant holds before, inside, and after the loop. The first condition is simply that the invariant I must evaluate to $true$ before the loop. Next, we need to assert that no matter what values the variables inside the loop have, the invariant and loop condition will hold whenever we go into the loop. Lastly, the invariant and the negated loop condition must hold when the loop terminates. We must also exchange all the occurrences in the currently accumulated weakest liberal precondition Q for these variables.

Additions to the standard rules. Most of the rules are the standard rules for computing the weakest liberal precondition. We have added two rules: the rule for `violate`, which will always give *false*, and the rule for assertions, which adds the assertion P to the accumulated condition Q, s.t. the new WLP is $P \wedge Q$.

Resolving undefined behaviour. As mentioned in the previous subsection, division and modulo by zero is undefined behavior. It is resolved by adding a condition when-

ever encountering a division or modulo operator, verifying that the operation is legal. For example, if we have an assignment $x = a\%b$, we will need the following precondition

$$\begin{aligned}&\forall y.((\neg(b = 0))\\ &\quad \wedge (b \neq 0 \Rightarrow y = a\%b \Rightarrow Q[x \leftarrow y]))\end{aligned}$$

which only resolves the assignment if b is not zero, and otherwise will always fail. Resolving division is similar. This modification is not explicitly written in the rules, but follows from the semantics of our *While* language.

2.3.2 Weakest Precondition

The weakest liberal precondition does not prove termination. However, by the variant in the while-loop introduced in Section 2.2.3, we can extend the *weakest liberal precondition* to *weakest precondition*, which will also ensure termination.
The structure for computing the WP for the constructs of the language is much like the one for computing the WLP, except for the structure of while loops, which is presented in Figure 4.

$$WP\left(\begin{array}{l}\texttt{while } e \texttt{ invariant } I\\ \quad \texttt{variant } v, \prec \texttt{ do } s\end{array}, Q\right) = I \wedge$$

$$\begin{aligned}&\forall x_1, ..., x_k, \xi,\\ &\quad (((e \wedge I \wedge \xi = v) \Rightarrow WP(s, I \wedge v \prec \xi))\\ &\qquad \wedge ((\neg e \wedge I) \Rightarrow Q))[w_i \leftarrow x_i]\end{aligned}$$

where $w_1, ..., w_k$ is the set of assigned variables in statement s and $x_1, ..., x_k, \xi$ are fresh logical variables.

Figure 4: The rule for computing the weakest precondition for while loops.

The difference between the computation of WP and WLP for while loops is the presence of the *variant*. When given a variant v for the loop, we assert that v decreases with each iteration, using the logical variable ξ to keep the old value of v to compare with.
For our example program `mult.ifc` we can compute the WP, as we have both a variant and an invariant in the while loop. By applying the WP rules on the example, we get the WP seen in Figure 5.
The interesting thing is how the while loop is resolved. First the *invariant* is checked in line 4, according to the first part of the *wp* rule. Next, the rule requires a universal quantifier over all the variables altered in the loop body, which happens in line 5. Inside the universal quantifier, we check that the loop invariant and condition hold when entering each iteration (lines 6-9) and that the invariant and the negated loop condition hold when the loop terminates (line 10). The variant is used in line 6, where the logical variable ξ is set to hold the value of the variant, and in line 9, where it is checked that the variant has decreased by comparing to the initial value of v stored in ξ.

```
1  ∀q, r. (q ≥ 0) ⇒
2    ∀res3. res3 = 0 ⇒
3      ∀$a. $a = q ⇒
4        (res3 = ($a − q) ∗ r ∧ q ≥ 0)
5         ∧∀q2, res2, ξ1.
6           (q2 > 0 ∧ res2 = ($a − q2) ∗ r ∧ q2 ≥ 0 ∧ ξ1 = q2) ⇒
7             ∀res1. res1 = res2 + r ⇒
8               ∀q1. q1 = q2 − 1 ⇒
9                 (res1 = ($a − q1) ∗ r ∧ q1 ≥ 0 ∧ 0 ≤ ξ1 ∧ q1 < ξ1)
10       ∧((q2 ≤ 0 ∧ res2 = ($a − q2) ∗ r ∧ q2 ≥ 0) ⇒ res2 = $a ∗ r)
11
```

Figure 5: Code for computing weakest precondition of `mult.ifc`

2.4 Verifying conditions with automated solvers

SMT-solvers The WP that we can generate will be a first order logic formula, which we want to validate. For this we use Satisfiability Modulo Theories (SMT). This is a modfication of the boolean satisfiability problem (SAT). Whilst SAT is limited to propositional logic and thus not expressive enough to reason about the verification conditions, SMT can be used to reason about first order logic, through a set of theories. A theory T is a set of sentences. A first order logic formula ϕ is satisfiable modulo this theory if there exists a model M such that $M \models_T \phi$. A verification condition for a program should be a valid formula and not just satisfiable, since it must hold for all models. However the problem of validating formula ϕ can be reduced to the problem of satisfying $\neg\phi$, since a formula must be valid if none of its negations are satisfiable.
For this project we need two theories to express the verification condition. Firstly we need the theory of Linear arithmetics, which provided reasoning about *aexpr*. Secondly, we need the theory of quantifiers. The SMT-solver of choice for this project is $Z3$[9], which uses the *SMT-Lib* standard, allows for both these theories, more specifically the theory is called LIA[1].
In practice the way SMT-solvers treat a first order logic formula is by using the theories to reduce the formula to some problem which can then be expressed as a SAT-formula. Hence, with a negation of a verification condition for a program we can feed it to a SMT-solver to find out if the condition is valid, implying the program is correct.

Why3. To ensure that our program is working as expected, we want to compare it to Why3, a well established tool for program verification. In Why3, theories can be built from existing theories, or from scratch, depending on what you want to prove. Most importantly Why3 allows us to define functions for which verification conditions can be generated and discharged to a variety of SMT solvers[**why3**]. One of these SMT-solvers are $Z3$, hence if we validate a formula in our language we should also be able to do so for an Why3 equivalent program. Hence it provides a well established target for verifying that our program works correctly. We compare our programs to equivalent Why3 programs. The semantics will be explained in Section 3.5.

3 Implementation

This section presents our implementation of an application for verification of programs, which we will refer to as *IFC*. The application consists of four parts, and each part carry

out a separate task for the program. This is done to make the code modular, such that each part does not explicitly depend on any other parts. The four parts of the application performs the following tasks:

1. **Parser:** Transforms an input program, written in *While*, to an Abstract Syntax Tree
2. **Interpreter:** Interprets a program, expressed by an AST, given an initial store
3. **Verification Condition Generator:** Computes the Verification Condition formula of a program, expressed by an AST, using predicate transformer semantics
4. **External SMT-solver API:** Run a verification condition formular through an external prover (Z3), to assert program correctness

Although each part can work seperately, they are connected in the application, as seen in Figure 6.

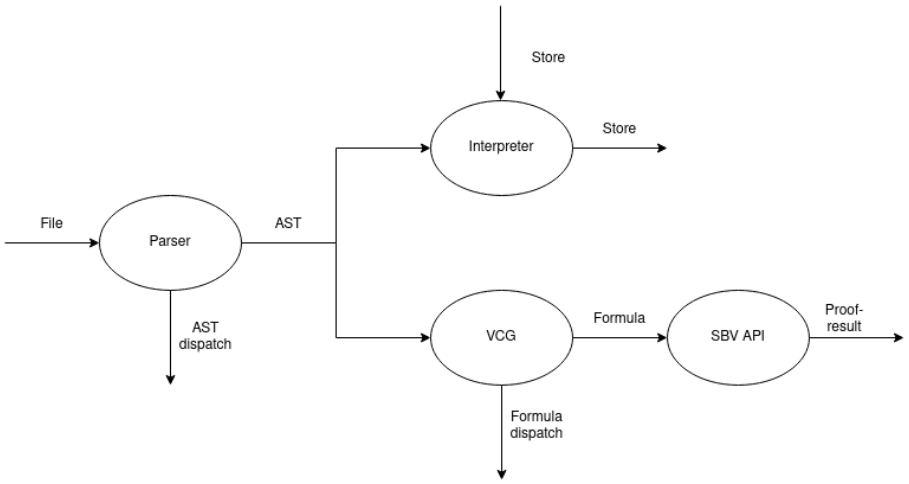

Figure 6: Application layout for IFC

Appendix A explains how to interact with the application. In the Section 3.1 to Section 3.4 we describe key points of the implementation, and Section 3.5 describes an interface which allows for more general proofs about programs.

3.1 Parser

The parser parses an input program, written in *While*, into an abstract syntax tree. The abstract syntax is very similar to the grammar in Table 1 and can be seen in Appendix B. Most of the implementation follows directly from the grammar, however, in the matter of resolving ambiguity, and introducing of ghost variables, we have made non-trivial design choices. These are presented in the following subsections.

3.1.1 Grammar ambiguity

The parser is built using the parser combinator library MegaParsec, which means we have to eliminate ambiguity in the grammar presented. We do so in two different

ways. For arithmetic operators and boolean operators, we make use of the expression parser defined in `Control.Monad.Combinators.Expr`. This makes handling operators, precedence and associativity easy. Furthermore, it allows for easy extension with new operators.
For parsing first order logic in assertions, we found the expression-parser unfit. One reason for this is that we want to allow for syntactic sugar, such as "∀ x y z.", which should desugar to "∀ x. ∀ y. ∀ z". Therefore, instead of using the expression-parser, we manually introduce precedence, according to conventions. That is, negation binds tightest, then conjuction, disjuction, quantifers, and lastly implication. Furthermore the grammar has been left-factorized. The resulting grammar can be seen in Figure 7.

$$
\begin{array}{lll}
\langle imp \rangle & ::= & \langle quant \rangle\ \langle imp' \rangle \\
\langle imp' \rangle & ::= & \text{'}\Rightarrow\text{'}\ \langle quant \rangle\ \langle imp' \rangle \mid \epsilon \\
\langle quant \rangle & ::= & \text{'}\forall\text{'}\ \langle vname \rangle\ \text{'.'}\ \langle quant \rangle \\
 & \mid & \text{'}\exists\text{'}\ \langle vname \rangle\ \text{'.'}\ \langle quant \rangle \\
 & \mid & \langle or \rangle \\
\langle or \rangle & ::= & \langle and \rangle\ \text{'}\vee\text{'}\ \langle or \rangle \mid \langle and \rangle \\
\langle and \rangle & ::= & \langle neg \rangle\ \text{'}\wedge\text{'}\ \langle and \rangle \mid \langle neg \rangle \\
\langle neg \rangle & ::= & \text{'}\sim\text{'}\ \langle factor \rangle \mid \langle factor \rangle \\
\langle factor \rangle & ::= & \langle bexp \rangle \mid \text{'('}\ \langle imp \rangle\ \text{')'}
\end{array}
$$

Figure 7: Modified grammar of *While*.

We also introduce some syntactic sugar in the grammar, such as "if c {s};" which will be desugared into "if c {s} else {skip};". This can easily be resolved in the parser by using the `option` parser-combinator.

For other parts of the grammar, which could have been syntactic sugar, such as implication and exist, we use the unsugared constructs. Although this introduces more code, the intent is to make the code easier to reason about in terms of the semantics. Furthermore it reads better in the output of the vc generator, and hence translate more directly into the predicate transformer semantics. This at least has made the development process easier. Ideally this could be alleviated by a more comprehensible pretty-printer, but at the moment, we settle for a slightly bigger AST.

3.1.2 Ghost variables

As previously mentioned in Section 2.2, it is semantically disallowed to use ghost variables anywhere in the program logic, except in assertions. Although this is a semantical matter, we handle this in the parser. If a ghost variable occurs in an invalid context, parsing simply fails, thus treating it as a syntactic issue. We do so by adding a reader monad transformer to the internal transformer type `parsecT` of `Megaparsec`.

```
type Parser = ParsecT Void String (ReaderT Bool Identity)
```

By changing the boolean value in this environment when entering an assertion context, we denote whether we are allowed to parse ghost variables at a certain point in the program. By only setting the boolean to *true* when in an assertion environment, we ensure that parsing of ghost variables will never happen in the program logic.

3.2 Interpreter

The interpreter follows directly from the semantics presented in Section 2.1. That is the intial store provided to the evalutor will be modified over the course of the program according to the semantics. We define the `Eval` type as follows such:

```
type STEnv = M.Map VName Integer
type Eval a = RWST () () STEnv (Either String) a
```

At the moment there is no use for neither `reader` nor `writer`, however when the language in the future is extended to have procedures the reader monad will be a natural choice for the scoping rules of said procedures. Likewise it is highly likely that the language would need to support some sort of output in the future. The store described in Section 2.1 is kept in the State `STEnv`. The store is simply a map from VNames to Integers. We want the store to be a State as the store after one monadic action should be chained with the next monadic action. This ensures that all variables are in scope for the rest of the program and hereby entails mutability. Duly note that ghost variables will also reside in this environment but will not be mutable or even callable, as previously explained. We make use of the error monad to resolve any runtime-errors that would arise, that is if a violation occurs, a ghost is assigned, a variable is used before it is defined, or if undefined behaviour arises such as division by 0. Semantically, the first error that occurs will be the return value of the computation.

Each type defined in the AST, `Stmt`, `FOL`, `AExpr`, `BExpr` is evaluated by different functions which all operate under the `Eval` monad, which allows for a clean and modular monadic compiler. They have the following types:

```
eval :: Stmt -> Eval ()
evalFOL :: FOL -> Eval Bool
evalBExpr :: BExpr -> Eval Bool
evalAExpr :: AExpr -> Eval Integer
```

From this we notice that all of the different constructs except for `Stmt`, will produce a value, whereas `Stmt` can only produce monadic actions, in terms of modifying the store. This monadic context allows us to translate the operational semantics almost directly into Haskell code. Figure 8 presents the code for evaluating statements.
The most complex argument for equivalence between the code and the semantics is the "while" construct. Figure 8, line 11-16, shows how we evaluate it. We first evaluate the invariant, to directly follow the semantical rules. Hence, when the invariant does not hold, we handle this case as abnormal termination, as per rule *while-i-false* from **??**, similarly to how we would do for a standard assertion. If the variant holds, but the loop-condition does not, we do nothing, as per rule *while-false*. Lastly if both the invariant and the loop-condition evaluates to *true*, we evaluate the body, and then we evaluate the while loop again. The other statements are simple and we treat them similarly to "while", directly in accordance with the semantics.

One important note about the interpreter is that we have no good way of checking assertions which includes quantifiers. The reasons is that we wanted to support arbitrary precision integers. This entails that we currently have no feasible way to check

```
eval :: Stmt -> Eval ()
eval (Seq s1 s2) = eval s1 >> eval s2
eval (GhostAss vname a) =
    get >>= maybe (update vname a) (const _e) . M.lookup vname
eval (Assign vname a) = update vname a
eval (If c s1 s2) =
    evalBExpr c >>= \c' -> if c' then eval s1 else eval s2
eval (Asst f) = evalFOL f >>= \case
  True  -> return ()
  False -> _e
eval w@(While c invs _var s) =
  evalFOL invs >>= \case
    False -> _e
    True -> evalBExpr c >>= \case
      True -> eval s >> eval w
      False -> return ()
eval Skip = return ()
eval Fail = _e
```

Figure 8: Evaluator for statements.[1]

such assertions. A potential solution would be to generate a symbolic reprensentation of the formula and try to satisfy it by using an external prover. A downside to this approach is that the interpreter would then also require external dependencies, and not be a standalone program anymore. All in all, this is unideal, and currently the approach is to ignore such assertions by considering them true. In **??** we describe a potential extension to *While*, which could help alleviate this problem. Non-quantified assertions are still evauluated as per the operational semantics.

3.3 Verification Condition Generator

The verification condition generator uses weakest precondition calculus, or weakest liberal precondition if a while loop has no specified variant, to construct the verification condition of a program. Like the approach in the interpreter we want to chain the actions and include a state and reader environment in the construction of the verification condition. This is done by using the following code:

```
type Counter = M.Map VName Count
type Env = M.Map VName VName
type WP a = StateT Counter (ReaderT Env (Either String)) a
```

The `Counter` state is used to give unique names to variables. The purpose of the reader environment is to resolve variable substitution in the formula generated from the weakest precondition calculus.

As described in Section 2.3, we use the quantified rule for assignment, when encountering assignments, to easily handle multiple occurences of the same variable. In our initial approach for implementing the VC generator, we tried to minimize the number of times we had to resolve condition Q, but our approach to this did not work. The working solution is a more naive approach. The next subsections presents both the initial attempt and the current approach.

3.3.1 First attempt

The initial approach tried to minimize the number of times we traversed condition Q. we tried to only resolve Q once, after the entire formula was built. This would allow resolving a sentence in just two passes, one over the imperative language AST and one over the structure of the formula. The approach was intended to build up a map as such, where `VName` is an identifier:

```
type Env1 = M.Map VName [VName]
```

Whenever encountering a variable we would add a unique identifier to its value-list along and modify Q by the weakest precondition for assignment. the initial pass over the AST of the program, would then be a partially resolved formula. The second traversal will be on this partly resolved formula. Whenever encountering a variable x, we would then look it up in the Map, and then replace x with the head of the list, since this must have been the last introduced bound variable. The values of the map can be seen as a sort of stacks. To resolve x with the appropriate bound variable, we would update the list each time a quantifier is encountered. The first time a quantifier is encountered we do nothing, as this simply means the first bound variable in the list is the one that should currently be used. Next time we encounter a quantification we would "pop" the first element, thus moving on to the next bound variable, and all occurrences of x from that point on will be substituted with the new bound variable, until encountering yet another quantification.
The problem with this approach is how our AST for First Order Logic formula is constructed. Consider the following example:

```
r := 5
r := r + 10
```

The first traversal would give the paritally resolved formula

$$\forall r \,.\, r = 5 \Rightarrow \forall r \,.\, r = r + 10 \Rightarrow true$$

and a tuple $(r, [r_1, r_2])$ for holding the list of bound variables for r. Now the result of the second traversal would then look like

$$\forall r_1 \,.\, r_1 = 5 \Rightarrow \forall r_2 \,.\, r_2 = r_2 + 10 \Rightarrow true$$

As can should be apparent that the last r in the formula should actually be substituted with r_1 and not by r_2 to make the formula correct according to the rules. But as r_1 has already been popped off the stack we substitute with r_2 instead, which is wrong. On the other hand, if we dont "pop" as soon as the second quantification is encountered we have no information on when to do so, as it is not necessarily the case that there appears an r in the equality before the implication. When this problem was dicovered we turned to the naive solution presented in the next subsection. It might be possible to include additional pattern matches such as `Forall x (AImp (Eq (Var y) r) rest)`, where r would then contain x_0, where $y_0 = y$. Here y should be substituted by x whereas y_0 should be substituted by the recently popped identifier. However we decided to go for the simpler approach.

3.3.2 Second attempt

The second and current approach is to resolve Q whenever we encounter an assignment. We generate the forall as such:

```
wp (Assign x a) q = do
    x' <- genVar x
    q' <- local (M.insert x x') $ resolveQ1 q
    return $ Forall x' (Cond (RBinary Eq (Var x') a) .=>. q')
```

First we make a new variable by generating a unique identifier, based on the State. Then we proceed to resolve q with the new environment, such that every occurence of variable x will be substituted by the newly generated variable stored in x'. The `Aexpr` which x evaluates to should not be resolved yet, as it could depend on variables not yet encounted. Ghost variables are easy to resolve, as they are immutable, and thus can be resolved right away.

The current version does not require the formula to be closed, but this is necessary for generating symbolic variables. To resolve this, we could simply do another iteration over the formula, checking that all program variables contain a `#`. However, we find that since the formula is intended to be fed to the next stage in the compiler, it is uneccesary to do so.

3.3.3 While loops with invariants and variants

The `while` statement has the most complicated rule in the weakest precondition calculus. The code for computing wp for a while loop is presented in Figure 9.

```
wlp (While b inv var s) q = do
  (fa, var', veq) <- maybe (return (id, ftrue, ftrue)) resolveVar var
  inv' <- fixFOL inv
  bs' <- fixBExpr b
  var'' <- maybe (return []) fixAExpr var
  w <- wlp s (inv ./\. var')
  fas <- findVars s []
  let inner = fa \$ reqs (((Cond b ./\. inv ./\. veq) .=>. w)
              ./\. ((anegate (Cond b) ./\. inv) .=>. q)) (inv' <> bs'
    <> var'')
  env <- ask
  let env' = foldr (\(x,y) a -> M.insert x y a) env fas
  inner' <- local (const env') \$ resolveQ inner
  let fas' = foldr (Forall . snd) inner' fas
  return \$ inv ./\. fas'
  where
    resolveVar :: Variant -> WP (FOL -> FOL, FOL, FOL)
    resolveVar v = do
      x <- genVar "variant"
      let cs' = Cond (bnegate (RBinary Greater (IntConst 0) (Var x)))
    ./\.
                Cond (RBinary Less v (Var x))
      return (Forall x, cs', Cond (RBinary Eq (Var x) v))
```

Figure 9: Weakest precondition for `while`

We attempt to make the code generic in terms existence of the variant, to eliminate code duplication. `ResolveVar` generate the conditions needed for the variant. If no variant is defined we produce a triple of identities. Is there a variant, we generate the equality $\xi = v$, along with the well-founded relation for unbounded integers. In the future we might have multiple ways of defining a well-founded relations, but this approach should translate well to other types. Line 3-5 will check if we need conditions

for Modulo and Division by 0, of which the approach is presented in Section 3.3.4. The rest of the code simply checks which variables are assigned in the body of the while loop, and generates a variable for each. And constructs the weakest precondition as per Figure 3 or Figure 4, depending on the presence of an invariant. Here it is important to note that the tool we provide does not itself distiguish between partial correctness and total correctness, this responsibility lies solely at the user.

3.3.4 Handling undefined behaviour in Verification Condition Generator

As described, we need to generate additional conditions whenever we encounter modulo and division. We do so by traversing arithmetic expressions to see if they might contain either modulo or division. This is quite a troublesome process as we need to check it multiple places and also need to traverse boolean conditions and assetions. When traversing the constructs we build up a list of conditions, which we need to apply. The lists, which we build will be of the form

```
[\x -> anegate (Cond (eq b (IntConst 0)))
   ./\. (Cond (bnegate (eq b (IntConst 0))) .=>. x), ...]
```

We can then fold these functions together on the construct that was traversed. We do not find this approach to be anywhere near ideal, and specifically because it generates formulas that are difficult to comprehend. The approach that *Why3* provides, where division is defined as a function which does this check, is much more comprehensible, but requires functions. Hence we find that this solution is reasonable for now, but should be replaced in the future.

3.4 Proof-assistant API

The proof-assistant API uses the SMT Based Verification library (SBV), which simplifies symbolic programming in Haskell. The library is quite generic and extensive compared to what we need. We mostly make use of the higher level functions, not utilising any internal functions.
Because the default type of SBV does not quite fit our needs, we instead use the provided transformer `SymbolicT`, to embed the `Except` monad. We want to do so as, when iterating over the formular, we might encounter a variable not yet defined and here fail gracefully, instead of throwing an error. Generating a `Predicate ~ Sym SBool` is relatively simple, since the formula generated in the previous stage is already a first order logic formula, making it straight forward to convert it into SBV's types. For the entire highlevel logic we resolve it as presented in Figure 10.
It is equally straight forward to resolve `bexpr` and `aexpr`. Ideally we would add a ReaderT to the transformer-stack to get rid of the explicit state. We have not been able to resolve the type for this, because of the following type constraint

```
forAll :: MProvable m a => [String] -> a -> SymbolicT m SBool
```

and because `m` must be an `ExtractIO`, which Reader and State does not implement, and which by the documentation cannot be implemented.
The predicate constructed by traversing the formula from the last stage is then be used as argument for the SBV function `prove`, which try to prove the predicate using Z3. If the program can be proved by the external SMT solver, the output will be `Q.E.D.`. If the formula is falsifiable, a caounter example is presented. For instance the output of the following program will obviously always be falsified.

```
type Sym a = SymbolicT (ExceptT String IO) a
type SymTable = M.Map VName SInteger

fToS :: FOL -> SymTable -> Sym SBool
fToS (Cond b) st = bToS b st
fToS (Forall x a) st = forAll [x] $ \(x'::SInteger) ->
  fToS a (M.insert x x' st)
fToS (Exists x a) st = forSome [x] $ \(x'::SInteger) ->
  fToS a (M.insert x x' st)
fToS (ANegate a) st = sNot <$> fToS a st
fToS (AConj a b) st = onlM2 (.&&) (`fToS` st) a b
fToS (ADisj a b) st = onlM2 (.||) (`fToS` st) a b
fToS (AImp a b) st = onlM2 (.=>) (`fToS` st) a b
```

Figure 10: Code for converting a first order logic formular into a symbolic bool.

```
violate;
```

whereas the multiplication program in Figure 2 is provable.

3.5 Interface for proofs

As described briefly in Section 2.1 we require that programs have a program header. This can also be seen from the example program presented earlier in this report. The header must look as follows:

```
vars: [ <variables> ]
requirements: { <preconditions> }
<!=_=!>
```

where
In the current implementation, there are no procedures in *While*, which makes it difficult to reason about the input of variables. Therefore we need the header to be able to generate and prove the weakest precondition for a program. The inspiration comes from how the *whyML* language defines procedures. Figure 11 is a whyML program equivalent to the `mult.ifc` program.
It is possible for why3 to generate a vector of input variables and then a precondition for each the conditions in `requires`, such that

$$\forall x_1, ..., x_n.\ requires \Rightarrow WP(body, ensures)$$

That is, whenever the *requires* holds, then the weakest precondition of the body should hold, where *ensures* states the postcondition. Notice that Figure 11 does not state any *ensures*, this more closely resembles our language. And since we dont have any return values, we dont really need *ensures*, since this might as well be an assertion in the actual program.

4 Assessment

In this section, we make an assessment of the correctness of our application for verification of *While* programs. We also assess the robustness and maintenance of the implementation, especially regarding any possible extension of the implementation. To support the assessment, we have designed a test suite consisting of both automated

```
module Mult

  use int.Int
  use ref.Refint

  let mult (&q : ref int) (r: int) : int
    requires { q >= 0 && r >= 0 }
    =
    let ref res = 0 in
    let ghost a = q in
    while q > 0 do
      invariant { res = (a - q) * r && q >= 0}
      variant { q }
      decr q;
      res += r
    done;
    assert { res = a * r };
    res
end
```

Figure 11: Why3 program equivalent to mult.ifc

tests, blackbox tests, and test programs. In Section 4.1, we present the experiments conducted as part of the testing strategy, and in Section 4.2 we assess the code based on the results of the tests, and perform a collective evaluation of our work.

4.1 Experiments

We have conducted a variety of tests to assess the implementation, consisting of both automatic tests and blackbox tests, and some manual experimentation using *Why3*. To thoroughly test our implementation, the tests follow this testing strategy:

- The *Parser* is tested using a blackbox approach, which should cover all aspects of the grammar presented in Table 1. Furthermore, we have a property based QuickCheck test to test for larger program constructs.
- The test suite includes property based testing to test the semantics of the *While* language, which compares the result of running the *Interpreter* on two semantically equivalent generated program constructs.
- We have attempt to use property based testing to compare interpretation of generated programs with the result of using our VC-generator coupled with *Z3* on the same program. This part is lacking since we want to generate some meaningful building blocks, which showed to be quite challenging and not critical part of this project.
- We use QuickCheck to generate random input for our example programs, which we use to run the programs and compare with the result of running the same computation in Haskell. This is to assert that the example programs does indeed do as we expect them to, telling us whether they should be provable or not.
- We test that a provable verification condition also results in a correct dynamic evaluation.

- To support the automated tests, we have done some analysis of a selection of programs and compared them to equivalent *Why3* programs.

A substantial part of the test suites and the experimentation is based on example programs, which can be found in Appendix D. Some of the examples are working programs, and some are designed to fail. Appendix C presents the programs (leftmost column), a short description of what they do (middle column), and the results of running them through the *Interpreter* and *VC generator* (rightmost column). We will analyse a selection of the examples in Section 4.1.3.
In Sections 4.1.1 to 4.1.3 we present the testing strategy more thoroughly. First we present the strategy of our blackbox *Parser* tests. Next we present our approach to generating automatic tests with property based testing. Finally we go into details with the example programs, and how they are designed to test the functionality of the implementation.

4.1.1 Blackbox testing

To assert that our implementation correctly parses input programs, we have designed a blackbox test suite for systematically testing each construct of the grammar. The strategy is to test parsing of smaller constructs such as variable names and expressions, and then combine them into larger constructs such as statements. Taking a systematic approach we hope to cover all aspects of the grammar. We have designed both positive and negative tests for the *Parser*, to assert that it behaves as intended when given both good and bad input.[2] Furthermore, we use blackbox tests to assert that the *Parser* handles ambiguities, associativity and precedence correctly. The parsing tests are still seriously lacking in quality (and far from complete), but we find the parsing to be fairly decently covered through the example programs.

4.1.2 Quickchecking instances

Other than the blackbox tests for the parser, our tests are mostly property based. With these automatic tests we attempt to verify that

- That the *Parser* correctly parses larger program constructs.
- That the *Interpreter* correctly follows the semantics of the *While* language.
- That there is consistency between the *VC-generator* and the *Interpreter*.

We build generators for the AST constructs, to enable generation of arbitrary input programs, and then use the generated input to test the properties specified above.

Generating input. Using the QuickCheck interface, we define instances of *Arbitrary* for the different AST types. When doing this, there are certain important considerations. Firstly, to ensure that the size of the generated expressions do not explode, we use *sized expressions* to control expansion, and to ensure that we get a good distribution of the various constructs we use *frequency* to choose between them. Secondly, we want the number of possible variables to be limited, such that the *Interpreter* will not fail too often, by using variables that have not yet been defined. This is done by limiting the options for variable names to be only single character string (this might still be

[2]The test-setup is heavily inspired by OnlineTA test by Andrzej Filinski from the course Advanced Programming.

too much). Thirdly, while-loops might not terminate, hence we want to define a small subset, or a skeleton, for while-loops that we can be sure terminates. This task proved quite challenging, since we need to construct while loops that are somewhat meaningful, in that they must terminate, and the invariant must be enough for the program to be validated. We have been able to generate some while-loops but they are not meaningful enough for the verifier, however, they can still be used in the semantic equivalence tests between different program constructs.

QuickCheck properties. With the possibility of generating input for the tests, we move on to finding meaningful properties to test. In this test suite we use property based testing for the following:

- **Parser tests.** To complement the blackbox tests for the parser, we want the following property to hold: *parse (prettyprint a) = a*, where a is an arbitrary AST. This is supposed to assert that the parser can handle a lot of different combinations of constructs, potentially finding bugs that would not have been discovered through our systematic blackbox testing of simple constructs.

- **Semantic equivalence.** To test whether the Interpreter corretly implements the semantics of the *While* language, we have tested certain equivalence properties. From studying the semantic system for *While*, we have designed equivalence properties according to the semantics. Examples of such equivalence properties are `if true then s1 else s2` $\sim$ `s1` and `while false do s` $\sim$ `skip`. These relations are directly related to the small-step semantics. Note that we have not presented these. The properties uses generators for generating suitable input variables and statements, i.e. the body of an if-statement is generated automatically, but the equivalence relation is defined manually.

- **Evaluating a program vs solving with VC generator and Z3.** To ensure consistency between the static and dynamic evaluation we wanted to generate some arbitrary programs and check for consistency. However, testing this automatically is quite a complex situation, since we need to be able to have a strong enough loop-invariant to prove the correctness of the program. It has come to our attention that the generated programs are not very useful, and currently we have not been succesful in implementing such a property-test. We will come back to this in the assessment in Section 4.2.

The above bulletpoints presents the intuition behind the QuickCheck testing of the implementation. A presentation of the test results and assessment of the code will be given in Section 4.2.

4.1.3 Dynamic execution compared to static proofs

Besides the QuickCheck and blackbox testing, we experimented with the example programs to check the quality of our implementation. In this subsection we present some of the example programs written for testing, and describe how and why they are interesting. In this process we relate the programs to equivalent to *Why3* programs, to see the expressiveness of our language. Next we set forth the experiments that we conduct using the programs as input to both the *Interpreter* and the *VC generator*. Finally we explain how we compare dynamic execution to static verification of the programs.

Example programs. In Appendix C we present an overview of some of the example programs written to test the implementation. The program `always_wrong` tests the `violate` construct, and simply checks whether the program correctly fails. The `skip` program tests the `skip` command, by asserting that the store is unchanged after executing the command. The programs `assign` and `max` are very simple too, and computes the result without use of while-loops, thus testing simple statement constructs.
The more interesting examples uses while-loops, in which the invariants and variants are important. The following present some of them, what they test, and why they are interesting.

- `mult.ifc`. The program takes as input two variables q, r and computes $q \times r$. This is the example that we use throughout the report, because it showcases a sequence of statements including assignments, assertions, use of ghost variables, and while-loops with both an invariant and a variant and still easy to follow. Thus the program tests a simple combination of statements, including a while loop that always terminates, and therefore the program should be provable under total correctness. An interesting thing about the program is to investigate whether the program is provable with the given invariant and variant. The code for `mult.ifc` is shown in Figure 2. Furthermore we made a modified program `mult3.ifc` which multiplies 3 variables by addition and hence tests the usage of nested while loops.

- `collatz.ifc`. This program calculates the collatz sequence of the input variable n. The while loop will run until $n = 1$. The idea behind this program is to have a while loop for which we cannot provide a variant. We can prove partial correctness of the program, but cannot prove total correctness. Ie., if we assume that the while loop terminates, we know that n is indeed 1 at loop termination, but we can never prove that the loop terminates. The program utilises assertions in a way that ensures that the postconditon only holds assuming the loop terminates. Therefore, a loop variant is necessary to determine whether the program is correct or not.

- `undef.ifc`. This program is an example of negative testing. It has three assignments, each testing different undefined behaviour:

  1. Testing that modulo by zero is undefined behaviour, and that the leftmost error is output.
  2. Testing that using an undeclared variable is undefined behaviour, and that the leftmost error is output.
  3. Testing that the leftmost error is output.

  Running the program shows how modulo and division by zero, along with using undeclared variables, is undefined behaviour. Furthermore, it tests that we always output the leftmost error.

  To actually see how the application handles the different errors, one must outcomment the previous lines of code, as the first error encountered will be output.

  This program also shows how use of undefined behaviour is considered a user error, thus trying to prove this program results in a falsifiable example, as undefined behaviour implies *false*.

- `mccarthy.ifc`. This program calculates the McCarthy 91 function. It takes as input an integer x and computes the result according to the formula

$$f(x) = \begin{cases} x - 10 & \text{if } x \geq 100 \\ f(f(x + 11)) & \text{otherwise} \end{cases}$$

  The result can also be described as:

$$f(x) = \begin{cases} 91 & \text{if } x \leq 100 \\ x - 10 & \text{otherwise} \end{cases}$$

  The interesting thing about this function is that although we can express the nested recursion as a while loop and this function indeed is correct. We have no way of expressing a valid variant and we cannot provide strong enough invariants. For us to be able to do this, we need to allow for pairs in the variant and to have recursive functions calls in assertions. Hence, we here a case of a problem which shows the relative completeness of Hoare logic, as we cannot prove our program although it is in fact correct. We can ensure so by a Why3 program.

- `isqrt.ifc & isqrt_fast` `isqrt` calculates the integer square root of a non-negative integer. We can both prove correctness of this. We further looked at an algorithm from Why3's verified programs gallery, which uses Newtons method for calculating this[]. Unfortunately we are not able to prove this, although we can write a semantically equivalent program. The reason for this is that why3 needs to some transformations of the subgoals for the proof to even validate it. Our solutions is not able any of such transformations.

Experiments with example programs. We have automated tests for testing that the *Interpreter* can correctly evaluate a program, and for testing that the programs that are provable using the VC generator will also evaluate to *true* in the *Interpreter*.
The first type of test, testing the evaluation of programs, is done by generating random input for the example programs, and then asserting that the result is in fact what we expect. For examle, when evaluating the `mult.ifc` program with two random values, the result should be equal to the result of multiplying the two values in Haskell. It should be noted that we use generators for generating meaningful input to the programs, to be able to test with all kinds of valid input.
The second type of test asserts that provably correct programs will evaluate correctly as well. This is tested by first feeding the programs to the VC generator coupled with Z3, and then to the *Interpreter* with random input. Given that the program is provable the *Interpreter* evaluates all assertions to *true*. Note that if we only show partial correctness, the *Interpreter* might run forever, so we only test for terminating instances. If the program is falsified, then the test will run the program with the falsifiable instance and assert that the *Interpreter* will terminate abnormally.
Another method that we use for verifying the correctness of our implementation is to check whether our application produces a result equivalent to that of *Why3*. Because *Why3* can use the same solver (*Z3*), and the same theories, we can create *Why3* programs equvalent to our test programs, and assert that they both either succeed in proving correctness, or fail with a counter example. Furthermore we have both positive and negative test programs, testing that our implementation both succeeds when it should, and fails when a program cannot be proven correct.

```
1  ∀q, r. (q ≥ 0 ∧ r ≥ 0) ⇒
2    ∀res3. res3 = 0 ⇒
3      ∀$a. $a = q ⇒
4        (res3 = ($a − q) ∗ r ∧ q ≥ 0)
5         ∧∀q2, res2, ξ1.
6           (q2 > 0 ∧ res2 = ($a − q2) ∗ r ∧ ξ1 = q2) ⇒
7             ∀res1. res1 = res2 + r ⇒
8               ∀q1. q1 = q2 − 1 ⇒
9                 (res1 = ($a − q1) ∗ r ∧ 0 ≤ ξ1 ∧ q1 < ξ1)
10       ∧((q2 ≤ 0 ∧ res2 = ($a − q2) ∗ r) ⇒ res2 = $a ∗ r)
```

Figure 12: Generated verification condition for the modified `mult` program.

Provable by VC generation ensures successful evaluation. As described, we value consistency highly, and it should always hold true that if we can prove total correctness of a program, then the dynamic evaluation should give the expected result. Once again it is important to note this will only hold for total correctness and not necessarily for partial correctness. Oppositely, correct dynamic evaluation does not imply provable correctness. The reason for this is that we might not have provided strong enough assertions to generate an appropriate verification condition, whilst the dynamic execution just needs all assertions to evaluate to *true*. This might be even more apparent, considering that quantifiers does not work correctly in the *Interpreter*. If for example we have a program that uses *true* as a loop invariant, this will hold in each iteration during the dynamic execution, but will probably not be enough to prove any postcondition statically.

Lets take a closer look at the multiplication example program from Figure 2. We have previously argued that the code is correct and can correctly be proved by Z3, but if we relax some of the assertions in the program, this will no longer be the case. Consider exchanging the loop-invariant `?{res = ($a - q) * r /\ q >= 0}` with the looser invariant `?{res = ($a - q) * r}`. Now Z3 will no longer be able to prove the correctness of the program. The generated formula looks as shown in Figure 12.

It becomes quite apparent that the invariant is no longer strong enough to prove the condition, since the restriction on q_2 is too weak. Z3 gives us a falsifiable example where $q_2 = -3$, $res_2 = 6$, $\$a = 0$ and $r = 2$. Using these values, the two first conjunctions in line 4 and lines 5-9 will evaluate to *true*. In the last term in line 10, the LHS of the implication will evaluate to *true* as $(-3 \leq 0 \wedge 6 = 3 * 2)$ is correct, but the RHS will evaluate to *false*, since $6 \neq 0 * 2$ is not correct. Thus the counter example does in fact falsify the verification condition.

But we know from testing that the program does in fact behave as intended, so how can the prover falsifies the formula? We can verify that our application acts correctly, by doing the same modification to the whyML program in Figure 11. Trying to prove the program correctness gives a falsifiable counter example, as expected. By this, we have confidence that our implementation works correctly, and requires the necessary loop invariants.

4.2 Evaluation

In this section we give a collective assessment of the project. We find that the two most important aspects of our application are correctness and maintainability. A non-correct

implementation is useless and high maintainability allows for easy extension/modification to build a more robust and complete solution. Overall we find that we satisfy these criteria to an acceptable degree.

4.2.1 Correctness

Overall we find that our implementation coincide with the specifications described and the goals we wanted to achieve with this project. The implementation deviates from the formal semantics in a few places, such as for abnormal termination, but this is done purely for convenience.
We find this based on our tests, experiments and so forth, presented in the previous section. All blackbox tests and QuickCheck tests pass, except for some time-outs that sometimes occur in the automatic tests. We suspect this has to do with the way we generate program constructs, and this could definitely be improved, also in regard of ensuring more meaningful generated input and minimising the amount of garbage that is generated. TODO: Test programs???
The fact that we can distinguish programs under partial correctness from programs under total correctness, and that the interpreter and static verification gives equivalent results, is a good indication that we have been successful in our implementation. With this said, we are not fully satisfied with the test-suite, and with more time the property based tests would have been stronger (and never time out), and we would have more unit-tests for the basic concepts of the different modules. Furthermore, we need to address the quantifiers in the interpreter. It is inferior that these are not handled appropriately. As stated, this stems from a combination of two events. Firstly, that we went with unbounded integers for the standard arithmetic type, which makes it impossible to go through all possible values, and secondly that we have treated the internal workings of SMT-solvers mostly as a blackbox and thus have not thouroughly investigated possible techniques for handling quantifiers. Although it would be a good idea to use one of the strategies from an SMT-solver, this is quite a substantial topic that we did not focus on in this project.

4.2.2 Completeness

We find that we have accomplished our goals for this project except for the parts mentioned in the correctness section. Furthermore it would have been good both for the implementation as well as our general understanding of provability of correctness to have delved more into the internals of SMT-solvers. As of now, our understanding is still basic.

4.2.3 Robustness

Overall we find that our testing has been valuable in insuring robustness of the code. As mentioned we have found peculiarities in the abnormalilites, which is now resolved. This assures that we can never prove anything which would result in abnormal termination in the interpreter. In this philospohy we gracefully handle run-time errors and such. At the same time we use *error* when we find that something is an implementation error. Thus we make a clear destinction between user errors and implementation errors.

4.2.4 Maintainability

As mentioned, this project is supposed to serve as a base, with the intention of extending it. Hence we have valued maintainability and extensibility highly. We have done so through the following abstractions. Firstly we are using monad-transformers to enforce seperation of concern. For the interpreter, we have defined the *Eval* type in terms of *RWST*, even though we currently only uses the *State*-transformer and underlying *Either* monad. We do so as we highly anticipate that the other monadic effects will be useful in the future, for example the *Write* monad could come in handy for printing. Furthermore, the multistage approach for generating verification conditions should allow for easy extension to other SMT-backends. In the current solution we only support *Z3*, but allowing other SBV supported solvers should be trivial. Because we generate the verification condition in a seperate state, it should not be too cumbersome to extend the implementation.

In terms of extending the language with new constructs we consider the following cases.

- If the construct can syntactically be reduced to one of the statements already defined, then only the parser would need change.
- New statements will require changes in the parser, interpreter and the verification condition generator, but will most likely not have to modify anything in the SBV API, since the (toplevel) first order logic is already fully defined. Adding statements should never have to modify the already defined constructs.
- New types, such as collections, would require additions in all modules, but should not add too much extra complexity compared to what new statements would.
- Adding a type-system would be a very useful addition, but will most likely be one of the biggest changes, as it would probably require additions to the constructors of the AST, or for an additional typed-AST, and in this case the code for the already working parts would have to be modified.

All in all, we find that the code is maintainable and extendable, but some extensions will require some overhaul of the code. As a supporting argument, when we found an error in the way we handled division and modulo, we were able to fix it rather smoothly, and it required fairly little additional code.

5 Discussion and conclusion

We have in this report presented the background for our project, which includes both the formal specification of the language, the axiomatic systems for Floyd-Hoare logic, and the predicate transformer semantics which we use to generate verification conditions for a program. We have presented our implementation and essential design choices of the program. Lastly we have argued why we find that our specific implementation meet all the goals presented in the introduction.

Nævn disse?

We have argued that our program is correctly able to statically prove correctness of certain programs and dynamically evaluate same program with a correct result. We have done so informally by tests and by a set of example programs, which shows examples of both partial correctness and total correctness of certain programs. We have further looked at how existing solutions, such as *Why3*, does verification conditions, by comparing *While* programs to *WhyML*. In the current state, the language is still sparse and quite restrictive. For the language to be used for more than just toying around, we propose three extensions, which we find essential to the language, in Section 5.1.

5.1 Future work

We describe three extensions that would make the programming language more useful. The first extension is additional constructs for the language, which would give more expressive power to programs. The next is including more different program types, to allow for more usability. The final one is the extension to include PER logic in the implementation. This is the extension we find the most interesting, since the main idea of this project was to allow for reasoning about Information Flow of programs.

5.1.1 Procedures and arrays

We propose that the language is extended to include procedures and arrays. This would enable us to make more interesting examples programs, for example it would allow for better assertions. Imagine having a procedure which determines the maximum value of two integers. Then we could use this in assertions about variables in another procedure. Most definitely this would provide a better experience using the language. In a similar manner we can create more interesting programs if we have arrays, or atleast some sort of generic-collection, which we can operate on.

5.1.2 Type system

A typesystem would also be a good extension for better usability. It would allow us for having different types stored in variables, and if coupled with the other extension arrays also, which could reduce a lot of repeated code. Furthermore additional types could eliminate a lot of "generic" assertions. An example of this would be a type *Nat* over the natural numbers, which would eliminate the necessity of checks such as $x \geq 0$. Such types would allow us to fix our faulty implementation of quantifiers in the interpreter as well, by allowing only certain types to be quantified. And although bruteforcing all values of a bound integer would still be extremely poor performace wise it would be a solution to our problem. However, bruteforcing should probably not be used, but the type system definitely opens up possibilities.

5.1.3 PER-logic for Information Flow Control

As mentioned multiple times, this project initially arose to make way for an implementation of the *Partial Equivalence Relation* logic for information flow control[]. The idea behind this is to prove properties of information flow, similarly to how we have been able to prove properties about simple *While* constructs. From our assessment we find that the project is ready, or at least close to ready, to try to include this program logic. However, its remains unbeknownst whether we can generate verification conditions for the logic.

5.2 Conclusion

References

[1] Clark Barrett, Pascal Fontaine, and Cesare Tinelli. *The Satisfiability Modulo Theories Library (SMT-LIB)*. www.SMT-LIB.org. 2016.

[2] Stephen A. Cook. "Soundness and Completeness of an Axiom System for Program Verification". In: *SIAM Journal on Computing* 7.1 (1978), pp. 70–90. DOI: 10.1137/0207005. eprint: https://doi.org/10.1137/0207005. URL: https://doi.org/10.1137/0207005.

[3] Edsger W. Dijkstra. "Guarded Commands, Nondeterminacy and Formal Derivation of Programs". In: *Commun. ACM* 18.8 (August 1975), pp. 453–457. ISSN: 0001-0782. DOI: 10.1145/360933.360975. URL: https://doi.org/10.1145/360933.360975.

[4] Andrzej Filinski, Thomas Jensen, and Ken Friis Larsen. "A PER Logic for Secure Information Flow". Private copy.

[5] Robert W. Floyd. "Assigning Meanings to Programs". In: *Proceedings of Symposium on Applied Mathematics* 19 (1967), pp. 19–32. URL: http://laser.cs.umass.edu/courses/cs521-621.Spr06/papers/Floyd.pdf.

[6] C. A. R. Hoare. "An Axiomatic Basis for Computer Programming". In: *Commun. ACM* 12.10 (October 1969), pp. 576–580. ISSN: 0001-0782. DOI: 10.1145/363235.363259. URL: https://doi.org/10.1145/363235.363259.

[7] Claude Marché. *Lecture notes in MPRI Course 2-36-1: Proof of Program*. 2013. URL: https://www.lri.fr/~marche/MPRI-2-36-1/2013/poly1.pdf.

[8] Hanne R. Nielson and Flemming Nielson. *Semantics with Applications: An Appetizer*. Undergraduate Topics in Computer Science. Springer, March 2007. ISBN: 1846286913. DOI: http://dx.doi.org/10.1007/978-1-84628-692-6. URL: http://dx.doi.org/10.1007/978-1-84628-692-6.

[9] *Z3 official repository*. https://github.com/Z3Prover/z3. Accessed: 2022-01-04.

A How To Use

The source-code for IFC, can be found at `https://github.com/Spatenheinz/IFC`. The application has been build using stack with LTS-18.24. Furthermore to run tests and use the discharging ability of verification conditions, $Z3$ needs to be install and on you $PATH$.

Project can be build using

```
stack build
```

To run the program use

```
stack exec -- PROGRAM
```

where PROGRAM denotes one of the following commands

```
IFC -e   program.ifc st  (parse program & interpret
                         with initial store st :: [(var,value)])
IFC -p   program.ifc     (prettyprint program)
IFC -ast program.ifc     (prints AST of program)
IFC -vc  program.ifc     (prints verification condition)
IFC -v   program.ifc     (run verification condition through Z3)
```

To run the test-suite do

```
stack test
```

B AST

We leave out smartconstuctors for simplifying printing of formulas. These can be seen in the code (in the GitHub repository).

```
module AST where

type Header = ([VName], Maybe FOL)

data BExpr = BoolConst Bool
  | Negate BExpr
  | BBinary BoolOp BExpr BExpr
  | RBinary ROp AExpr AExpr
  deriving (Show, Eq)

data BoolOp = Conj | Disj
  deriving (Show, Eq)

data ROp = Less | Eq | Greater
  deriving (Show, Read, Eq)

data AExpr = Var VName
  | Ghost VName
  | IntConst Integer
  | Neg AExpr
  | ABinary ArithOp AExpr AExpr
  deriving (Show, Eq)

data ArithOp = Add | Sub | Mul | Div | Mod
  deriving (Eq, Show, Enum)

data Stmt = Seq Stmt Stmt
  | GhostAss VName AExpr
  | Assign VName AExpr
  | If BExpr Stmt Stmt
  | Asst FOL
  | While BExpr FOL (Maybe Variant) Stmt
  | Skip
  | Fail
  deriving (Show, Eq)

type Variant = AExpr

data FOL = Cond BExpr
  | Forall VName FOL
  | Exists VName FOL
  | ANegate FOL
  | AConj FOL FOL
  | ADisj FOL FOL
  | AImp FOL FOL
  deriving (Show, Eq)

type VName = String
```

C Table of test programs

program	description	results
always_wrong	This program is simply a `violate` statement and should thus always fail.	The interpreter evaluates to `false`, and the VC generator gives a falsifiable verification condition.
assign	Assigns values to two variables, and asserts that they get assigned correctly.	The interpreter gives the expected result, and the VC generation gives a solvable verification condition.
collatz	This program calculates the length of the collatz-conjecture until the repeating pattern $1, 4, 2, 1, ...$	The interpreter gives the expected result, but only partial correctness can be proved, as we cannot find an appropriate variant.
div	Takes as input two variables a, b and computes the euclidean division of a by b.	The interpreter gives the expected result, and the VC generator generates a provable verification condition.
div_in_cond	Tests the use of division by zero in conditions.	The interpreter correctly gives a division by zero error message, but the VC generator does not falsify. This seems to be due to a problem with the way we append assertions about errors in e.g. while loops.
fac	Takes as input an integer r and computes $r!$.	Our assertion language is not expressive enough to give an appropriate loop invariant and postcondition, thus the program cannot be proven, and evaluating it through the interpreter makes little sense.
fakesum	A sum program that has an incorrect postcondition (wrongly subtracts 1 from the result).	Both the interpreter and VC generator behaves as expected.
false_mult	A multiplication program that has an invariant that is too weak.	The interpreter evaluates to the correct result and verifies all assertions, as we would expect, but the VC generator cannot prove the correctness because of a too weak invariant, as is also expected.
fib	Takes as input an integer a, and computes the a'th Fibonacci number.	The evaulator gives the expected result, but the VC generator falsifies the computed verification condition.

infinity	Takes no input and runs forever. The postcondition asserts false, which will hold as the negated loop condition is *false*, and thereby everything holds.	The VC generator correctly outputs QED. For obvious reasons the interpreter cannot evaluate the program.
isqrt	Takes as input an positive integer x and computes the arithmetic squareroot of x.	The interpreter gives the expected result, and the VC generator computes a provable verification condition.
isqrt_fast	Another way of computing the arithmetic squareroot.	The interpreter gives the expected result, but the VC generator cannot prove the program, because one needs program transformations that our assertion language cannot express.
isqrt_sub	A third way of computing the arithmetic squareroot, using only subtraction.	The interpreter gives the expected result, and the VC generator computes a provable verification condition.
max	Takes as input two variables x, y and finds the maximum of the two.	The interpreter gives the expected result, and the VC generator computes a provable verification condition.
mccarthy	Takes as input an integer n and if $n \leq 100$ outputs 91, and otherwise outputs $n - 10$.	The assertion language is not expressive enough to express an appropriate loop variant or invariant, so the program cannot be proven. The interpreter finds the correct result and verifies all assertions.
mod0	Tests that modulo by zero is not allowed.	The interpreter gives the expected result, and the VC generator computes a falsifiable verification condition.
mult	Takes as input two variables q, r and multiplies them.	The interpreter gives the expected result, and the VC generator computes a provable verification condition.
mult3	Takes as input three variables q, r, s and computes $q \times r \times s$. This is done with a nested while loop.	The interpreter gives the expected result, and the VC generator computes a provable verification condition, although quite slowly (because of the naive nature of the program).

skip	Asserts that the store is unchanged when executing a `skip` command.	The interpreter gives the expected result, and the VC generator computes a provable verification condition.
sum	Takes as input a positive integer n and sums the integers from 1 to n.	The interpreter gives the expected result, and the VC generator computes a provable verification condition.
undef	Tests the various undefined behaviour: division and modulo by zero, and usage of undeclared variables.	Correctly falsifies the program, both in the VC generator and the interpreter.
var	Tests that we cannot prove correctness of a program that never terminates.	If there is no variant, our application succeeds in proving partial correctness, ass the postcondition is then simply *true*. However, when giving a variant the program is falsified, as the loop never terminates, and thus total correctness cannot be proven.

D Example programs

D.1 always_wrong.ifc

```
vars: []
requirements: {}
<!=_=!>
violate;
```

D.2 assign.ifc

```
-- yea we can assign stuff
vars: [i]
requirements: {}
<!=_=!>
a := i;
n := 10 + a;
#{n = i + 10};
```

D.3 collatz.ifc

```
-- Collatz terminates but we cannot show it.... yay for partial
vars: [n]
requirements: {n > 0}
<!=_=!>
k := 420;
while (n /= 1) ?{ n > 0}{
      if (n % 2 = 0) {
         n := n / 2;
      } else {
         n := 3 * n + 1;
      };
};
if n = 1 {
k := 42;
};
#{ k = 42};
```

D.4 div.ifc

```
vars: [a,b]
requirements: {0 <= a /\ 0 < b}
<!=_=!>
q := 0;
r := a;
while r >= b
?{ a = q * b + r /\ 0 <= r}
!{r}
{
r := r - b;
q := q + 1;
};
#{q * b + r = a /\ 0 <= r /\ r < b};
```

D.5 div_in_cond.ifc

```
-- this fails since we have division by 0.
vars: []
requirements: {}
```

```
<!=_=!>
while (10 / 0 < 20) ?{ false } {
skip;
};
```

D.6 fac.ifc

```
-- Can we express this? i dont think so
vars: [r]
requirements: {r > 0}
<!=_=!>
q := 1;
while r > 0 ?{r >= 0} !{r} {
  q := q * r;
  r := r - 1;
};
```

D.7 fakesum.ifc

```
-- Sum but -1 in post condition clearly not true.
vars: [n]
requirements: {n > 0}
<!=_=!>
sum := 0;
m := 0;
while (m <= n) ?{ sum = ((m-1) * (m)) / 2}; ?{n - m >= -1} !{n-m} {
      sum := sum + m;
      m := m + 1;
};
#{ sum = (n* (n + 1)) / 2 - 1};
```

D.8 false_mult.ifc

```
-- This has a too weak invariant
vars: [q,r]
requirements: {q >= 0}
<!=_=!>
res := 0;
$a := q;
while (q > 0) ?{res = ($a - q) * r} !{q} {
      res := res + r;
      q := q - 1;
};
#{res = $a * r};
```

D.9 fib.ifc

```
-- FIB, we cannot express this in assertions. Interpreter works tho.
vars: [a]
requirements: {a > 0}
<!=_=!>
fib1 := 0;
res := 1;
tmp  := 0;
while a > 1
?{ a > 0 } !{a} {
     tmp := res;
     res := res + fib1;
     fib1 := tmp;
     a := a - 1;
};
```

D.10 infinity.ifc

```
-- DONT PUT THIS IN TESTS
vars: []
requirements: {}
<!=_=!>
while true ?{true} {
skip;
};
#{false};
```

D.11 isqrt.ifc

```
-- The most basic isqrt
vars: [x]
requirements: {x >= 0}
<!=_=!>
count := 0;
sum := 1;
while sum <= x
?{x >= count * count};
?{sum = (count + 1) * (count + 1)}
!{x - count}
{
  count := count + 1;
  sum := sum + (2 * count + 1);
};
res := count;
#{res >= 0 /\ res * res <= x /\ x < (res + 1) * (res + 1)};
```

D.12 isqrt_fast.ifc

```
-- Logarithmic runtime (but you know assertions are slow)
vars: [x]
requirements: { x >= 0}
<!=_=!>
if x = 0 {
  res := 0;
} else {
   if x <= 3 {
     res := 1;
   } else {
     y := x;
     z := (1 + x) / 2;
     while z < y
     ?{z > 0};
     ?{y > 0};
     ?{z = (x / y + y) / 2};
     ?{x < (y + 1) * (y + 1)};
     ?{x < (z + 1) * (z + 1)}
     !{ y }
     {
        y := z;
        z := (x / z + z) / 2;
        a := x / y;
        #{ (x < a * y + y
         => a + y <= 2 * z + 1
         => (a + y + 1) * (a + y + 1) <= (2 * z + 2) * (2 * z + 2)
         => 4 * ((z + 1) * (z + 1) - x) = (2 * z + 2) * (2 * z + 2) - 4
     * x
```

```
            /\ (2 * z + 2) * (2 * z + 2) - 4 * x >= (a + y + 1) * (a + y + 1) - 4 * x
            /\ (a + y + 1) * (a + y + 1) - 4 * x > (a + y + 1) * (a + y + 1) - 4 * (a * y + y)
            /\ (a + y + 1) * (a + y + 1) - 4 * (a * y + y) = (a + 1 - y) * (a + 1 - y)
            /\ (a + 1 - y) * (a + 1 - y) >= 0)
        => x < (z + 1) * (z + 1) };
    };
    #{ y <= x / y => y * y <= (x / y) * y};
    res := y;
    };
};
#{res >= 0 /\ res * res <= x /\ x < (res + 1) * (res + 1)};
```

D.13 isqrt_sub.ifc

```
-- This is quite clever, does isqrt by subtraction
vars: [x]
requirements: {x >= 0}
<!=_=!>
b := 5;
a := b * x;
while (a >= b)
?{x >= (b / 10) * (b / 10)};
?{a = 5 * (x - (b / 10) * (b / 10))};
?{b = 5 + (b / 10) * 10};
?{b / 10 >= 0}
!{a} {
     a := a - b;
     b := b + 10;
};
res := b / 10;
#{res >= 0 /\ res * res <= x /\ x < (res + 1) * (res + 1)};
```

D.14 max.ifc

```
-- simple max program. MAX SPEED MAX FIDELITY
vars: [x,y]
requirements: {}
<!=_=!>
$a := x;
if x < y { x := y; };
#{(x = $a /\ $a > y) \/  x = y };
```

D.15 mccarthy.ifc

```
-- Fun recursive function as while loop. However assertion language not expressive enough
vars: [x]
requirements: {}
<!=_=!>
n := x;
c := 1;
while c > 0
?{c >= 0};
?{ (x > 100 => n >= 91) /\ (x <= 100 => n <= 111) }
{
    if n > 100 {
      n := n - 10;
      c := c - 1;
```

```
    } else {
      n := n + 11;
      c := c + 1;
    };
};
#{ (x <= 100 => n = 91) /\ (x > 100 => n = x - 10) };
```

D.16 mod0.ifc

```
-- since we modulo by 0 its false
vars: []
requirements: {}
<!=_=!>
$b := -15 % (25 + ((-9) - 16));
```

D.17 mult.ifc

```
-- Multiplies two integers
vars: [q,r]
requirements: {q >= 0}
<!=_=!>
res := 0;
$a := q;
while (q > 0)
?{res = ($a - q) * r /\ q >= 0}
!{q}
{
      res := res + r;
      q := q - 1;
};
#{res = $a * r };
```

D.18 mult3.ifc

```
-- We can also nest while loops thats kinda cool brah
vars: [q,r,s]
requirements: {q >= 0 /\ r >= 0 /\ s >= 0}
<!=_=!>
res := 0;
$a := q;
while (q > 0)
?{res = ($a - q) * r * s /\ q >= 0}
!{q}
{
      res2 := 0;
      i := r;
      while (i > 0) ?{res2 = (r - i) * s /\ i >= 0} {
        res2 := res2 + s;
        i := i - 1;
      };
      res := res + res2;
      q := q - 1;
};
#{res = $a * r * s};
```

D.19 skip.ifc

```
-- Checks that skip does nothing
vars: [a]
requirements: {}
```

```
<!=_=!>
$a := a;
skip;
#{ a = $a };
```

D.20 sum.ifc

```
-- Sum the first n numbers
vars: [n]
requirements: {n > 0}
<!=_=!>
sum := 0;
m := 0;
while (m <= n) ?{ sum = ((m-1) * (m)) / 2}; ?{n - m >= -1} !{n-m} {
      sum := sum + m;
      m := m + 1;
};
#{ sum = (n* (n + 1)) / 2};
```

D.21 undef.ifc

```
-- Leftmost thing is the abnormal
vars: []
requirements: {}
<!=_=!>
-- a := (10 % 0 + b) + (10 / 0);
-- a := (10 % (0 + b)) + (10 / 0);
c := (1 % (10 % 10 % 1)) / 0;
```

D.22 var.ifc

```
-- This program will QED if the variant is not present (obviously the
    loop never terminates), fails ow.
vars: [n]
requirements: {n > 0}
<!=_=!>
while (true)
?{true}
!{n}
{
  n := n - 1;
};
```